



University of Ljubljana
Faculty of Computer and
Information Science



GPU Optimized Math Routines in the **Stan** Math Library

Rok Češnovar, Davor Sluga, Jure Demšar, Steve Bronder, Erik Štrumbelj

MAIN **GOAL**

Faster model inference
for Stan users ...

... in a seamless and cost-
effective way.



TALK **OUTLINE**

- ▶ Motivation: GP regression.
- ▶ Parallelization & GPUs.
- ▶ Stan + OpenCL.
- ▶ Paralellizing the Cholesky decomposition.
- ▶ Challenges & Roadmap.



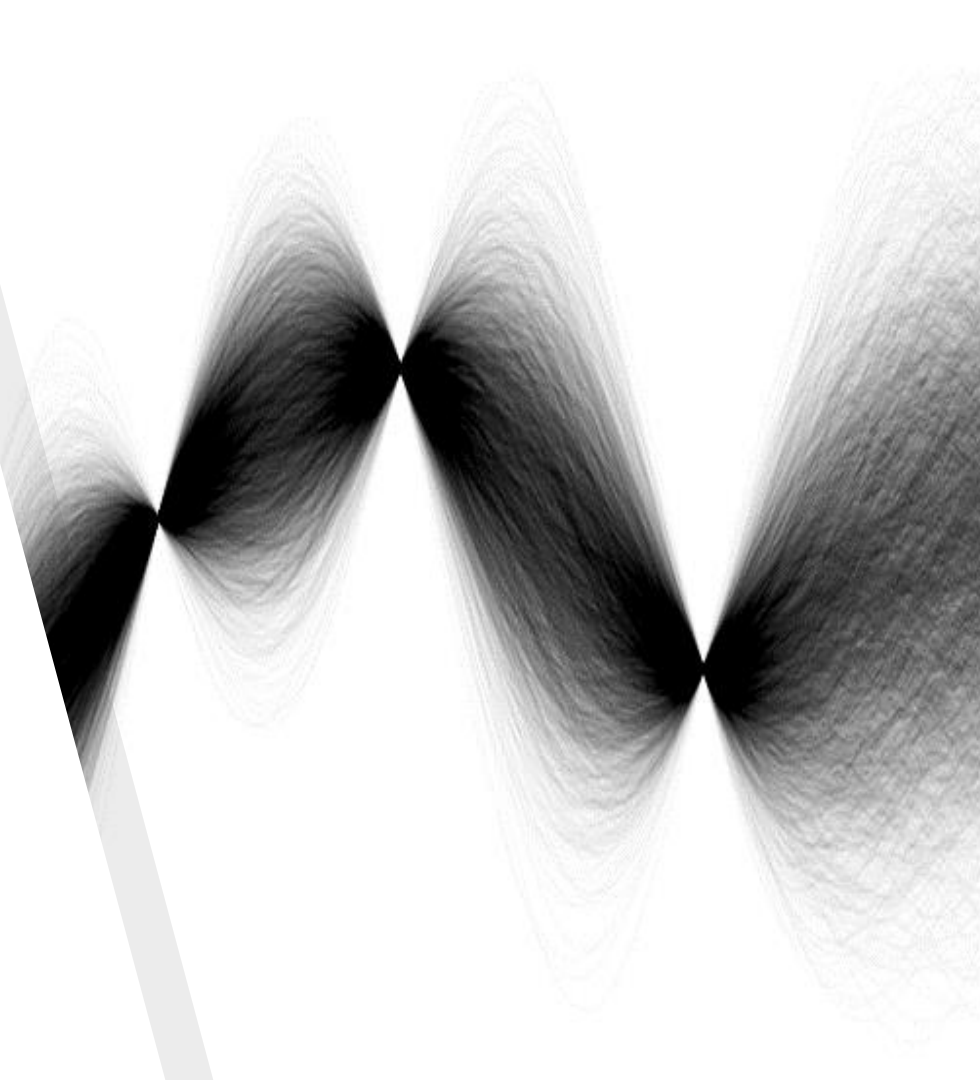
GP REGRESSION

Gaussian Processes are very useful but ...
... computation scales unfavorably - $O(n^3)$.



What can we do?

- ▶ Approximate inference.
- ▶ Run the computation on a better CPU.
- ▶ Run the computation in parallel.



PARALLELIZATION 101

Break the problem up into smaller pieces and compute them in parallel.

The pieces must be:

- ▶ Small enough to keep all individual processing units busy most of the time.
- ▶ Large enough to avoid too much overhead with breaking them up and putting them back together.

Maximum speedup limited by the parts we can't parallelize.



WHY GPU^s?

Properties of GPUs:

- ▶ Everyone has a GPU.
- ▶ Massive parallelism – thousands of „cores“ (best performance/cost ratio).
- ▶ Optimized for vector and matrix problems.
- ▶ Faster data transfers compared to clusters.
- ▶ Energy efficient.



Stan

+ OpenCL



OpenCL

OpenCL

- ▶ Parallel framework CPUs, GPUs, DSPs, FPGAs,...
- ▶ Special functions (kernels) are executed by N threads on target devices.
- ▶ APIs for C, C++, Python, Julia...
- ▶ Open standard maintained by the Khronos Group.

Typical example:

- ▶ Copy input data to the OpenCL device,
- ▶ parallel execution of special functions on the OpenCL device and
- ▶ copy the results back to the host.



OpenCL

What's an OpenCL **Context**?

It's like a scheduler:

- ▶ Manages the devices, queues, platforms, memory alloc, etc.
- ▶ Devices: GPUs and CPUs.
- ▶ Platforms: Implementations of OpenCL (Khronos's OpenCL vs. Intel's OpenCL).
- ▶ The context has a 'program' object that manages kernels for devices and platforms.



OpenCL

OpenCL Context

- ▶ We only want one context to exist so it's stored as a singleton.
- ▶ Access the context through an adapter API.

IE: In the adapter class `openc1_context`

```
// Return the stan program's context
inline cl::Context& context() {
    return openc1_context_base::getInstance().context_;
}
```

...

```
// developers access
auto ctx = openc1_context.context()
```



Making **Kernels**

- ▶ We want to make it simple for users to add and use kernels.
- ▶ Reworked design with Sean Talts and Rok.

Making a kernel:

A kernel that only needs to set the global work size.

Name of kernel

```
const global_range_kernel<cl::Buffer, cl::Buffer, int, int> copy("copy", copy_kernel_code);
```

Argument Types kernel accepts

String literal holding the kernel code

Accessing Matrices on the Device

Developers move Stan matrices over to
`matrix_cl` matrices:

```
matrix_d d1  
matrix_cl d1_cl(d1)
```

Users operate on these like Stan matrices:

```
matrix_cl d3_cl = d1_cl + d2_cl  
matrix_cl d4_cl = cholesky_decompose(d1_cl)
```



OpenCL


First GPU-optimization available to Stan users:

Cholesky Decomposition

Cholesky

Decomposition

- ▶ Our first bottleneck target in Stan.
- ▶ Almost no naive parallelisms in the basic algorithm.
- ▶ No real speedup on the GPU.
- ▶ Blocked-cholesky is computationally more complex but GPUs are made for fast matrix multiplications.



Derivative of the Cholesky Decomposition

GPU implementation of [Murray 2016]:

- ▶ Eigen version already in Stan Math,
- ▶ largest bottlenecks are matrix multiplication and lower triangular inverse.


As a consequence - GPU implementations of

- ▶ matrix multiplication,
- ▶ lower triangular inverse,
- ▶ +, -, transpose, partition.

```
function chol_blocked_rev(L,  $\bar{A}$ )  
  # If at input  $\text{tril}(\bar{A}) = \bar{L}$ , at output  $\text{tril}(\bar{A}) = \text{tril}(\bar{\Sigma})$ , where  $\Sigma = LL^\top$ .  
  for  $k = N$  to no less than 1 in steps of  $-N_b$ :  
     $j \leftarrow \max(1, k - N_b + 1)$   
     $R, D, B, C = \text{level3partition}(L, j, k)$   
     $\bar{R}, \bar{D}, \bar{B}, \bar{C} = \text{level3partition}(\bar{A}, j, k)$   
     $\bar{C} \leftarrow \bar{C}D^{-1}$   
     $\bar{B} \leftarrow \bar{B} - \bar{C}R$   
     $\bar{D} \leftarrow \bar{D} - \text{tril}(\bar{C}^\top C)$   
     $\bar{D} \leftarrow \text{chol\_unblocked\_rev}(D, \bar{D})$   
     $\bar{R} \leftarrow \bar{R} - \bar{C}^\top B - (\bar{D} + \bar{D}^\top)R$   
  return  $\bar{A}$ 
```

Lower triangular inverse

GPU Basic forward substitution algorithm not suitable for GPUs.

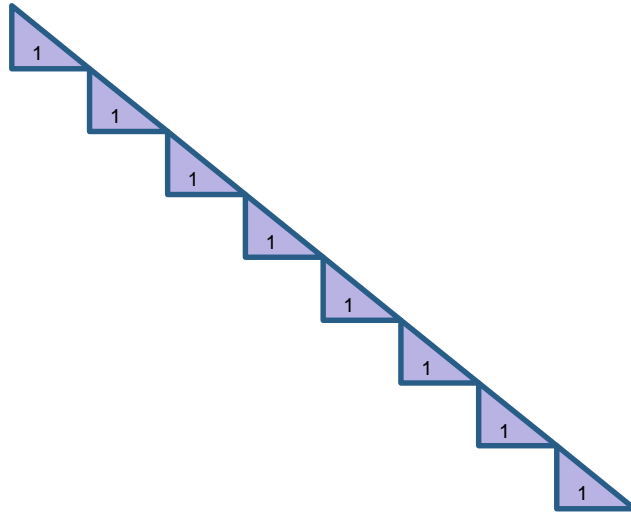


$$C_1 = A_1^{-1}$$

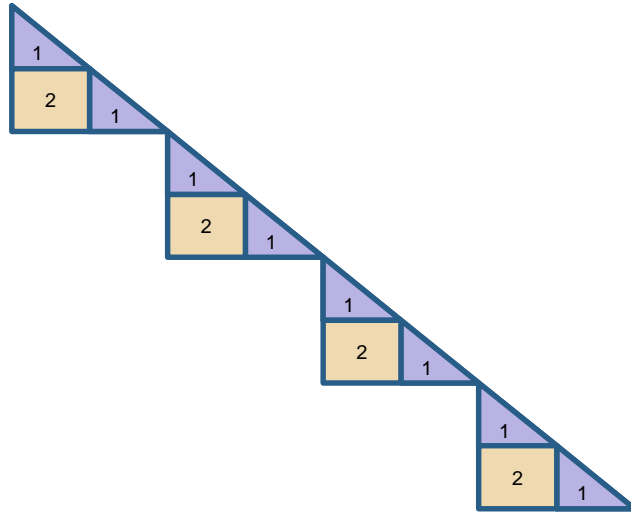
$$C_2 = A_2^{-1}$$

$$C_3 = -C_2 A_3 C_1$$


Lower triangular inverse



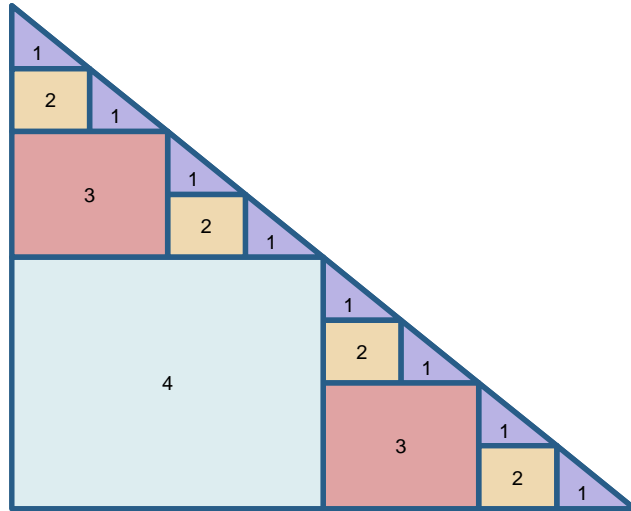
Lower triangular inverse





Lower triangular inverse




Lower triangular inverse



RESULTS (simple 1D GP regression on N points)



Issues & Roadmap



Issue: Data tranfers

Currently, speedups are limited by data transfers:

- ▶ Copying data to the GPU costs us some time.
- ▶ Most functions scale linearly and can't justify this cost for every call.
- ▶ Even is such functions represent 10% of the computation, it drastically limits the maximum speedup.

Example:

```
cov_exp_quad(x1, magnitude_1, length_scale_1)
+ cov_exp_quad(x1, magnitude_2, length_scale_2)
+ gp_periodic_cov(x1, magnitude_3, length_scale_3_1, 7)
.* cov_exp_quad(x1, 1.0, length_scale_3_2)
+ gp_periodic_cov(x1, magnitude_4, length_scale_4_1, 365.25)
.* cov_exp_quad(x1, 1.0, length_scale_4_2)
+ gp_dot_prod_cov(I_s, magnitude_5_1)
+ gp_dot_prod_cov(I_ss, magnitude_5_2)
+ gp_dot_prod_cov(I_ws, magnitude_5_3)
+ diag_matrix(rep_vector(jitter, N2));`
```

Multiple Devices and OOM Algorithms

- ▶ OpenCL can run on CPUs and GPUs, so why not use both?
- ▶ We know some problems are too small to send to the GPU, so use OpenCL on the CPU.
- ▶ Needs a 'smart' load balancer that can look at a 'job' and decide where to send it.
- ▶ Current algorithms are limited by GPU DRAM.
- ▶ We should be able to configure the algorithms to work in a chunking fashion.



From **single routines** to (almost) **entire models**

Idea: move the bulk of the log-posterior computation (including „hand-made“ gradients) to the GPU.

Promising results on some models (~100x speedup).

Currently in the works:

- ▶ GLM (linear, logistic, Poisson, NB regression...). For example:

```
bernoulli_logit_glm_lpdf(y | X, beta, alpha);
```

- ▶ Roadblock: Data transfers!

